


Schema Evolution Patterns

Alex Rasmussen

alex@bitsondisk.com



1. *Geospiza magnirostris*.
3. *Geospiza parvula*.

2. *Geospiza fortis*.
4. *Certhidea olivacea*.

John Gould (14.Sep.1804 - 3.Feb.1881) [Public domain]

Hi, I'm Alex!



**LA-based
Data Engineering
Consultant**



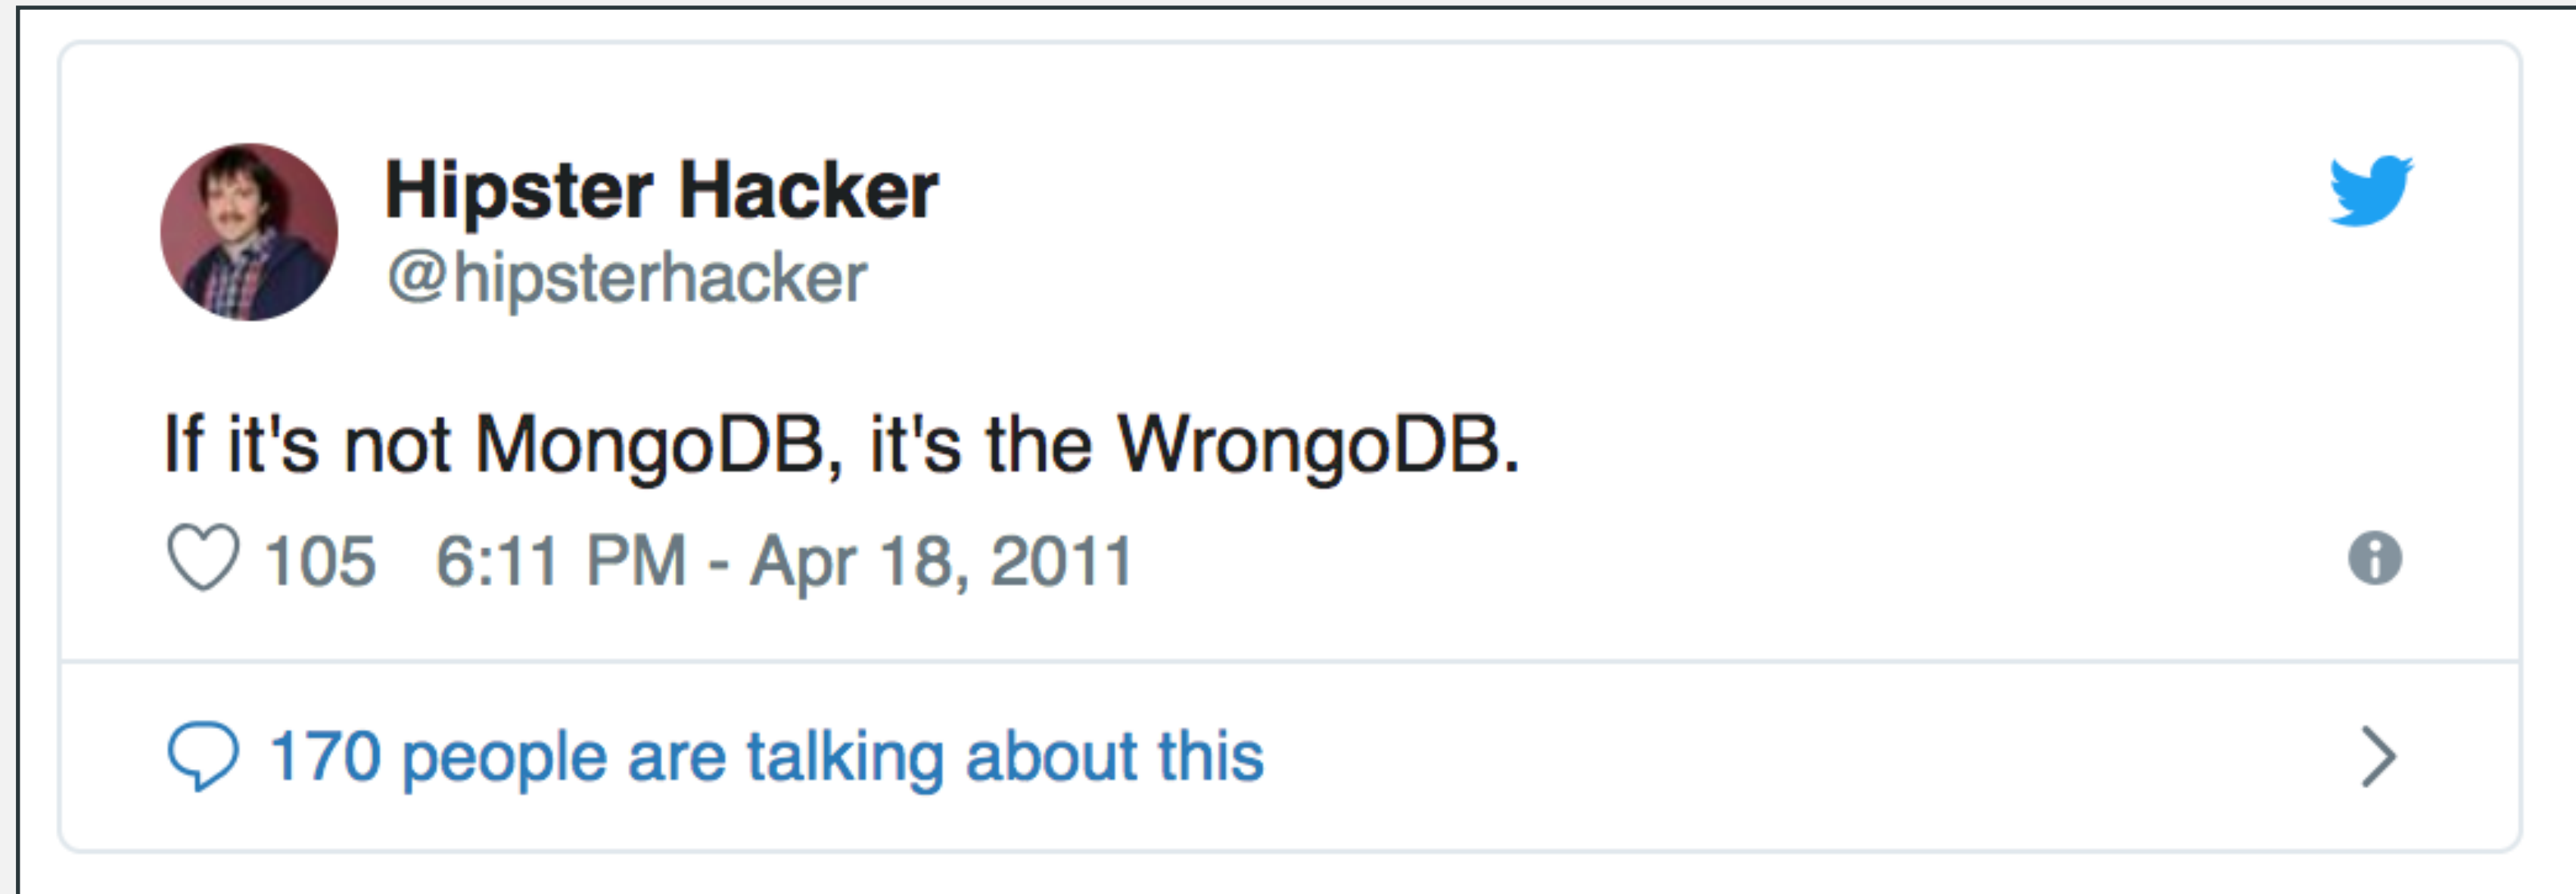
BITS ON DISK

<https://www.bitsondisk.com/>

**Twitter, GitHub, etc: @alexras
alex@bitsondisk.com**

Data Migrations Are Hard

- **How should I migrate? When?**
- **Will the old code still work?**
- **Can I migrate without downtime?**
- **Can I roll back? How? What if I can't?**



Should I just use MongoDB?
How do I decide?

API Versioning Is Hard

- **When should I bump versions?**
- **Will old callers break?**
- **Should I just switch to GraphQL?**

**THESE HARD
PROBLEMS
ARE THE SAME
PROBLEM!**

Schema Evolution

- The structure of your data **will** change
- What does that mean for the data?
- How should our systems respond?
- How does this impact **how we build**?

Goals of This Talk

- **Mental toolkit for reasoning about schema evolution that applies across technologies and domains**
- **How this looks in practice**

References: <https://github.com/bitsondisk/velocity-sj-2019>

1. WHAT ARE SCHEMAS?

2. SCHEMA COMPATIBILITY

3. CROSSING COMPATIBILITY GAPS

4. WRAPPING UP / TAKEAWAYS

1. WHAT ARE SCHEMAS?

2. SCHEMA COMPATIBILITY

3. CROSSING COMPATIBILITY GAPS

4. WRAPPING UP / TAKEAWAYS

Schemas

- From the Greek *skhēma*, meaning “form” or “figure”. The data’s **shape**.

```
Person {  
    name: string,  
    age: integer,  
    favorite_color: enum(RED, BLUE, ...),  
    . . .  
}
```

What/Where Is the Schema?

- **Explicit:** SQL tables, Thrift, Avro*
- **Implicit:** CSV, Excel, JSON
- **Separate:** JSON Schema, schema registries,
an ex-employee's brain

When Are Schemas Enforced?

- On **Write**: e.g. in RDBMS tables
- On **Read**: enforce after reading
- On **Need**: dynamic, late-binding
- Can enforce **one** schema, or **many**

1. WHAT ARE SCHEMAS?

2. SCHEMA COMPATIBILITY


3. CROSSING COMPATIBILITY GAPS

4. WRAPPING UP / TAKEAWAYS

Schema Compatibility


- If two schemas are **compatible**, evolving from one schema to another can be done **automatically** on read*
- Readers can be oblivious to schema change**
- Two directions: **backwards** and **forwards**

Backwards-Compatible



**Data written with old schema
readable by clients with new schema**

Forwards-Compatible



Data written with new schema
readable by clients with old schema

Add a Field With a Default

```
name: string,  
age: integer,  
is_admin: boolean (default: false)
```

```
name: "Bob Jones",  
age: 42
```



```
name: "Bob Jones",  
age: 42,  
is_admin: false
```

Backwards: substitute default in older data

Forwards: ignore field in newer data

Remove a Field With a Default

```
name: string,  
age: integer,  
is_admin: boolean (default: false)  
pto_days_left: int (default: 0)
```

```
name: "Alice Smith",  
age: 29,  
is_admin: true  
pto_days_left: 16
```



```
name: "Alice Smith",  
age: 29  
is_admin: true
```

Backwards: ignore field in older data

Forwards: substitute default in newer data

Other Types of Changes

- **Without defaults:**
 - **Adding a field breaks backwards-compatibility**
 - **Removing a field breaks forwards-compatibility**
- **Field type changes: typecasting rules apply(?)**
- **Renaming and reordering: it depends**

In Practice - Protocol Buffers

```
message Person {  
    required string name = 1;  
    required int32 age = 2;  
    optional bool is_admin = 3  
        [default = false];  
}
```

- **Field numbers can't change or be reused**
- **In proto3, no required or optional**

In Practice - Avro

```
record Person {  
  string name;  
  int age;  
  boolean is_admin = false;  
}
```

- **Reorders are OK, but not renames**
- **Compatibility determined by rules**

In Practice - GraphQL

```
type Person {  
  name: String!,  
  age: Int!,  
  is_admin: Boolean  
}
```

- **Every field is nullable by default**
- **Only return queried fields (bounded enforcement)**
- **Hide deprecated fields; deletes tricky**

In Practice - Wix

- **Adding fields**

- **Non-indexed fields:** JSON blob column (schema-on-read!)
- **Indexed fields:** new table + join

- **Removing fields**

- **Don't. Just stop using them (similar to GraphQL)**

Recap

- **Compatibility makes evolution easier**
- **Changes can be backwards-compatible, forwards-compatible, both, or neither**
- **Ease-of-compatibility drives many decisions behind “rules” of format design**

1. WHAT ARE SCHEMAS?

2. SCHEMA COMPATIBILITY

3. CROSSING COMPATIBILITY GAPS

4. WRAPPING UP / TAKEAWAYS

Crossing Compatibility Gaps

- Not all schema changes are **compatible**
- Not all incompatibilities are **simple**
- Not all compatible changes are **practical**
- How do we deal with that?

Complex Changes

```
name: string,  
first_name: string,  
last_name: string,  
age: integer,  
is_admin: boolean (default: false)
```

- **Not obvious how to split**
- **Client code changes (likely) required**


Impractical Changes

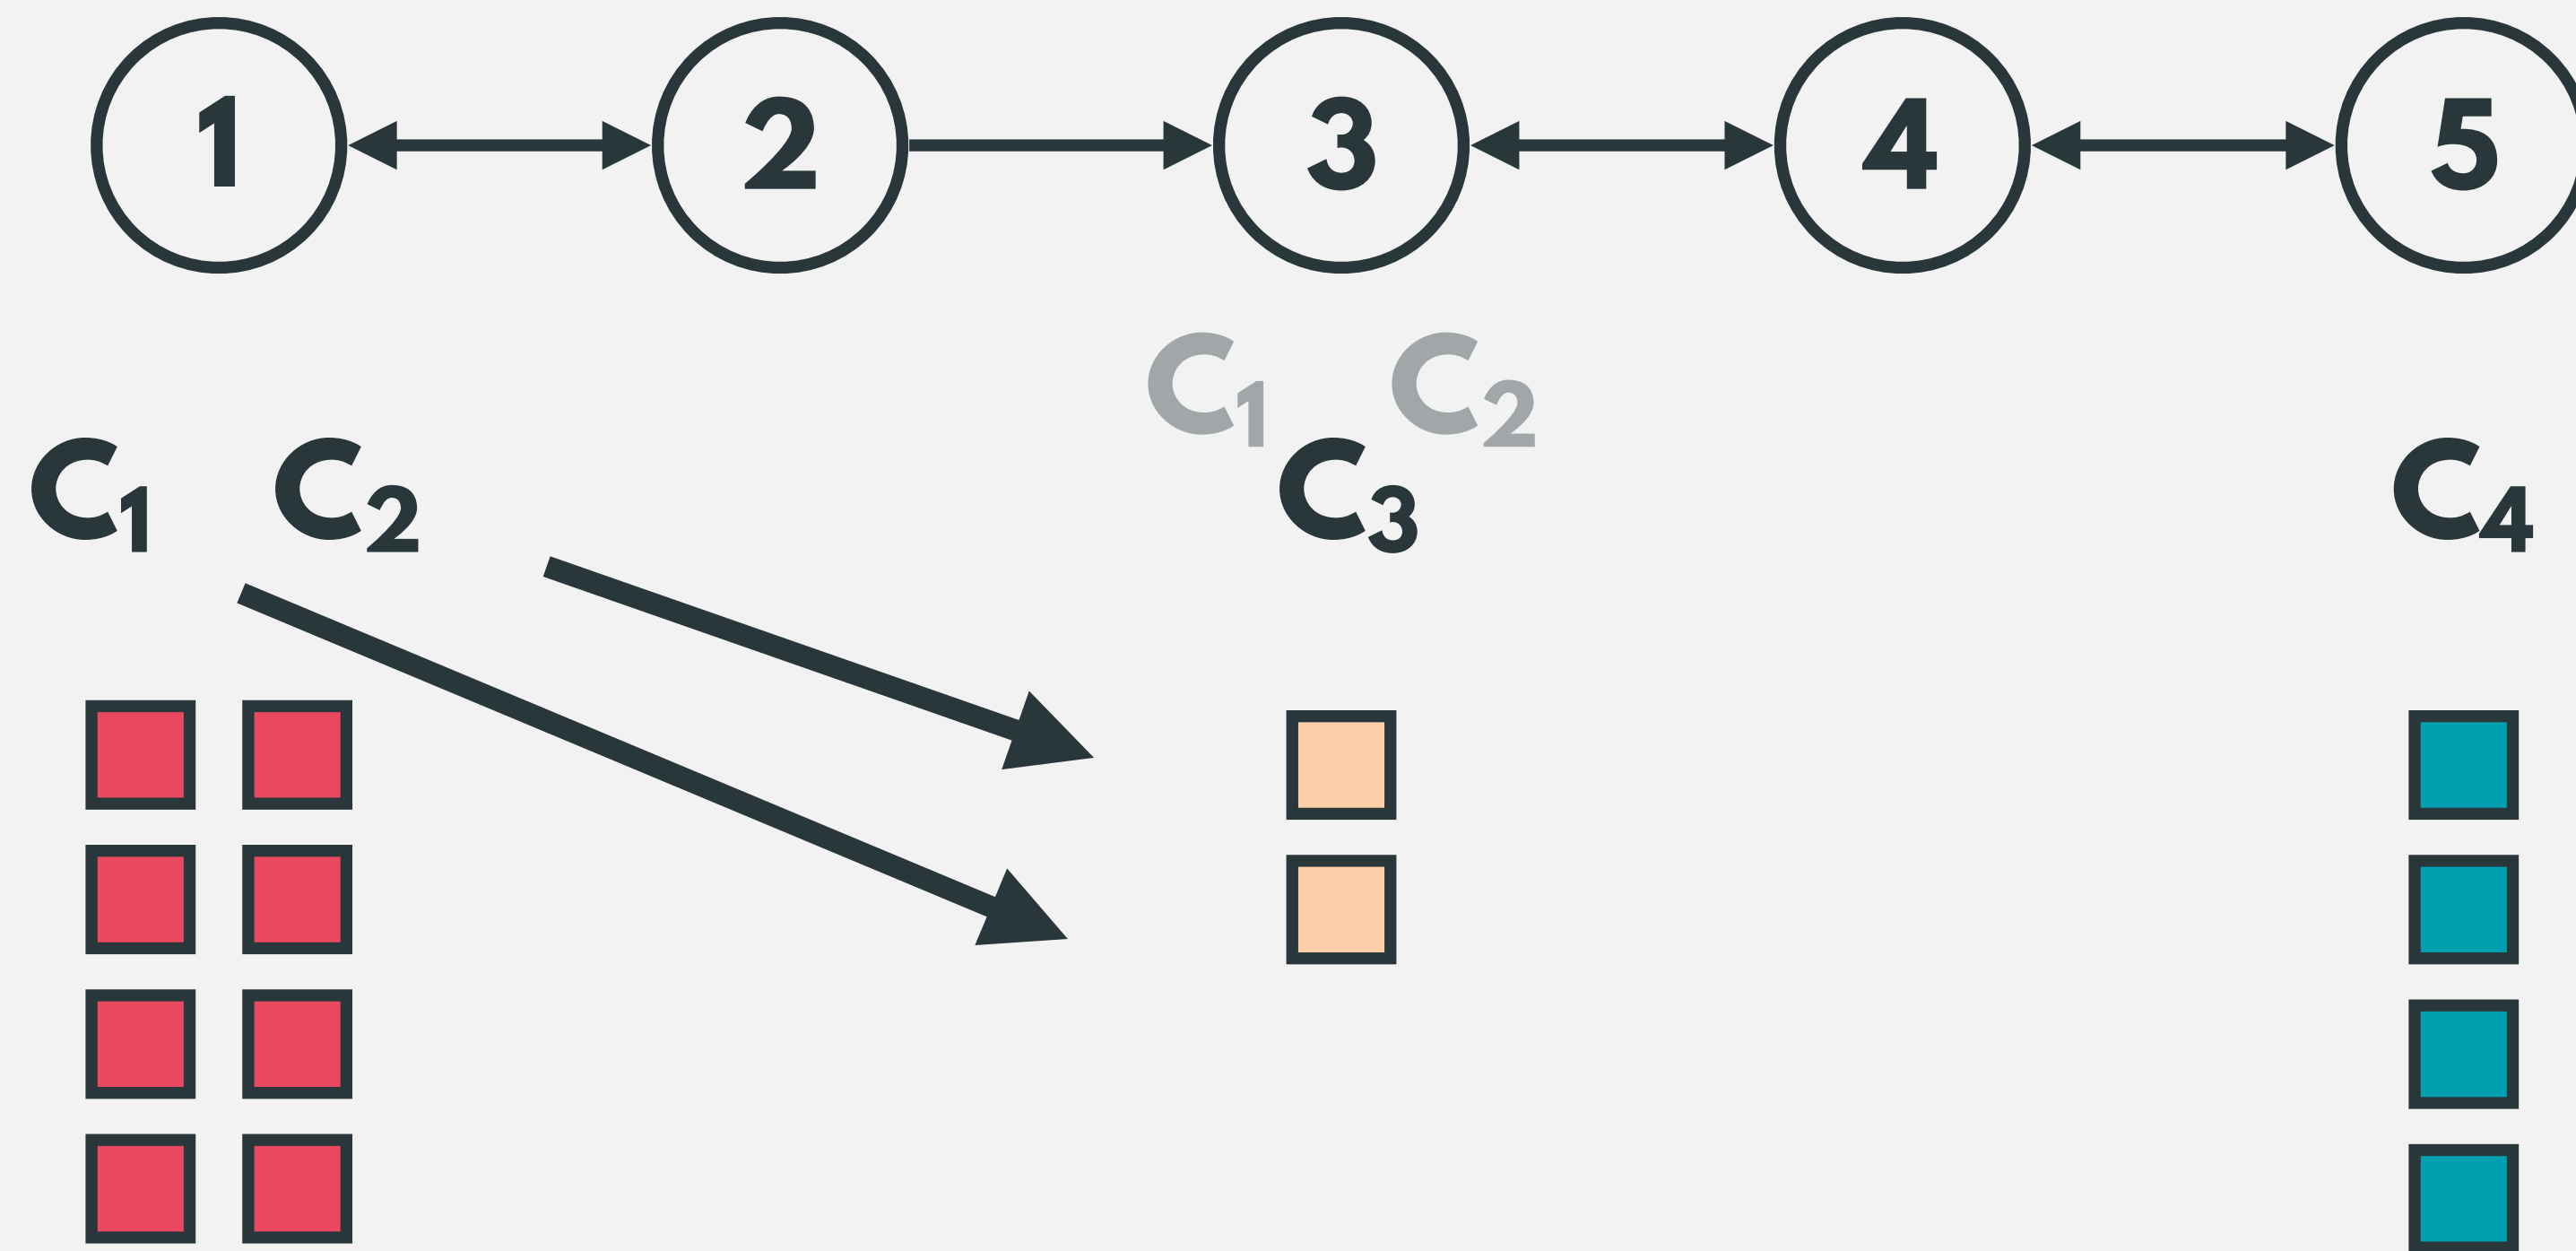
- e.g. Adding a column in *MySQL* requires locking/copying the table
 - **Days to weeks** not unheard of

Crossing Compatibility Gaps

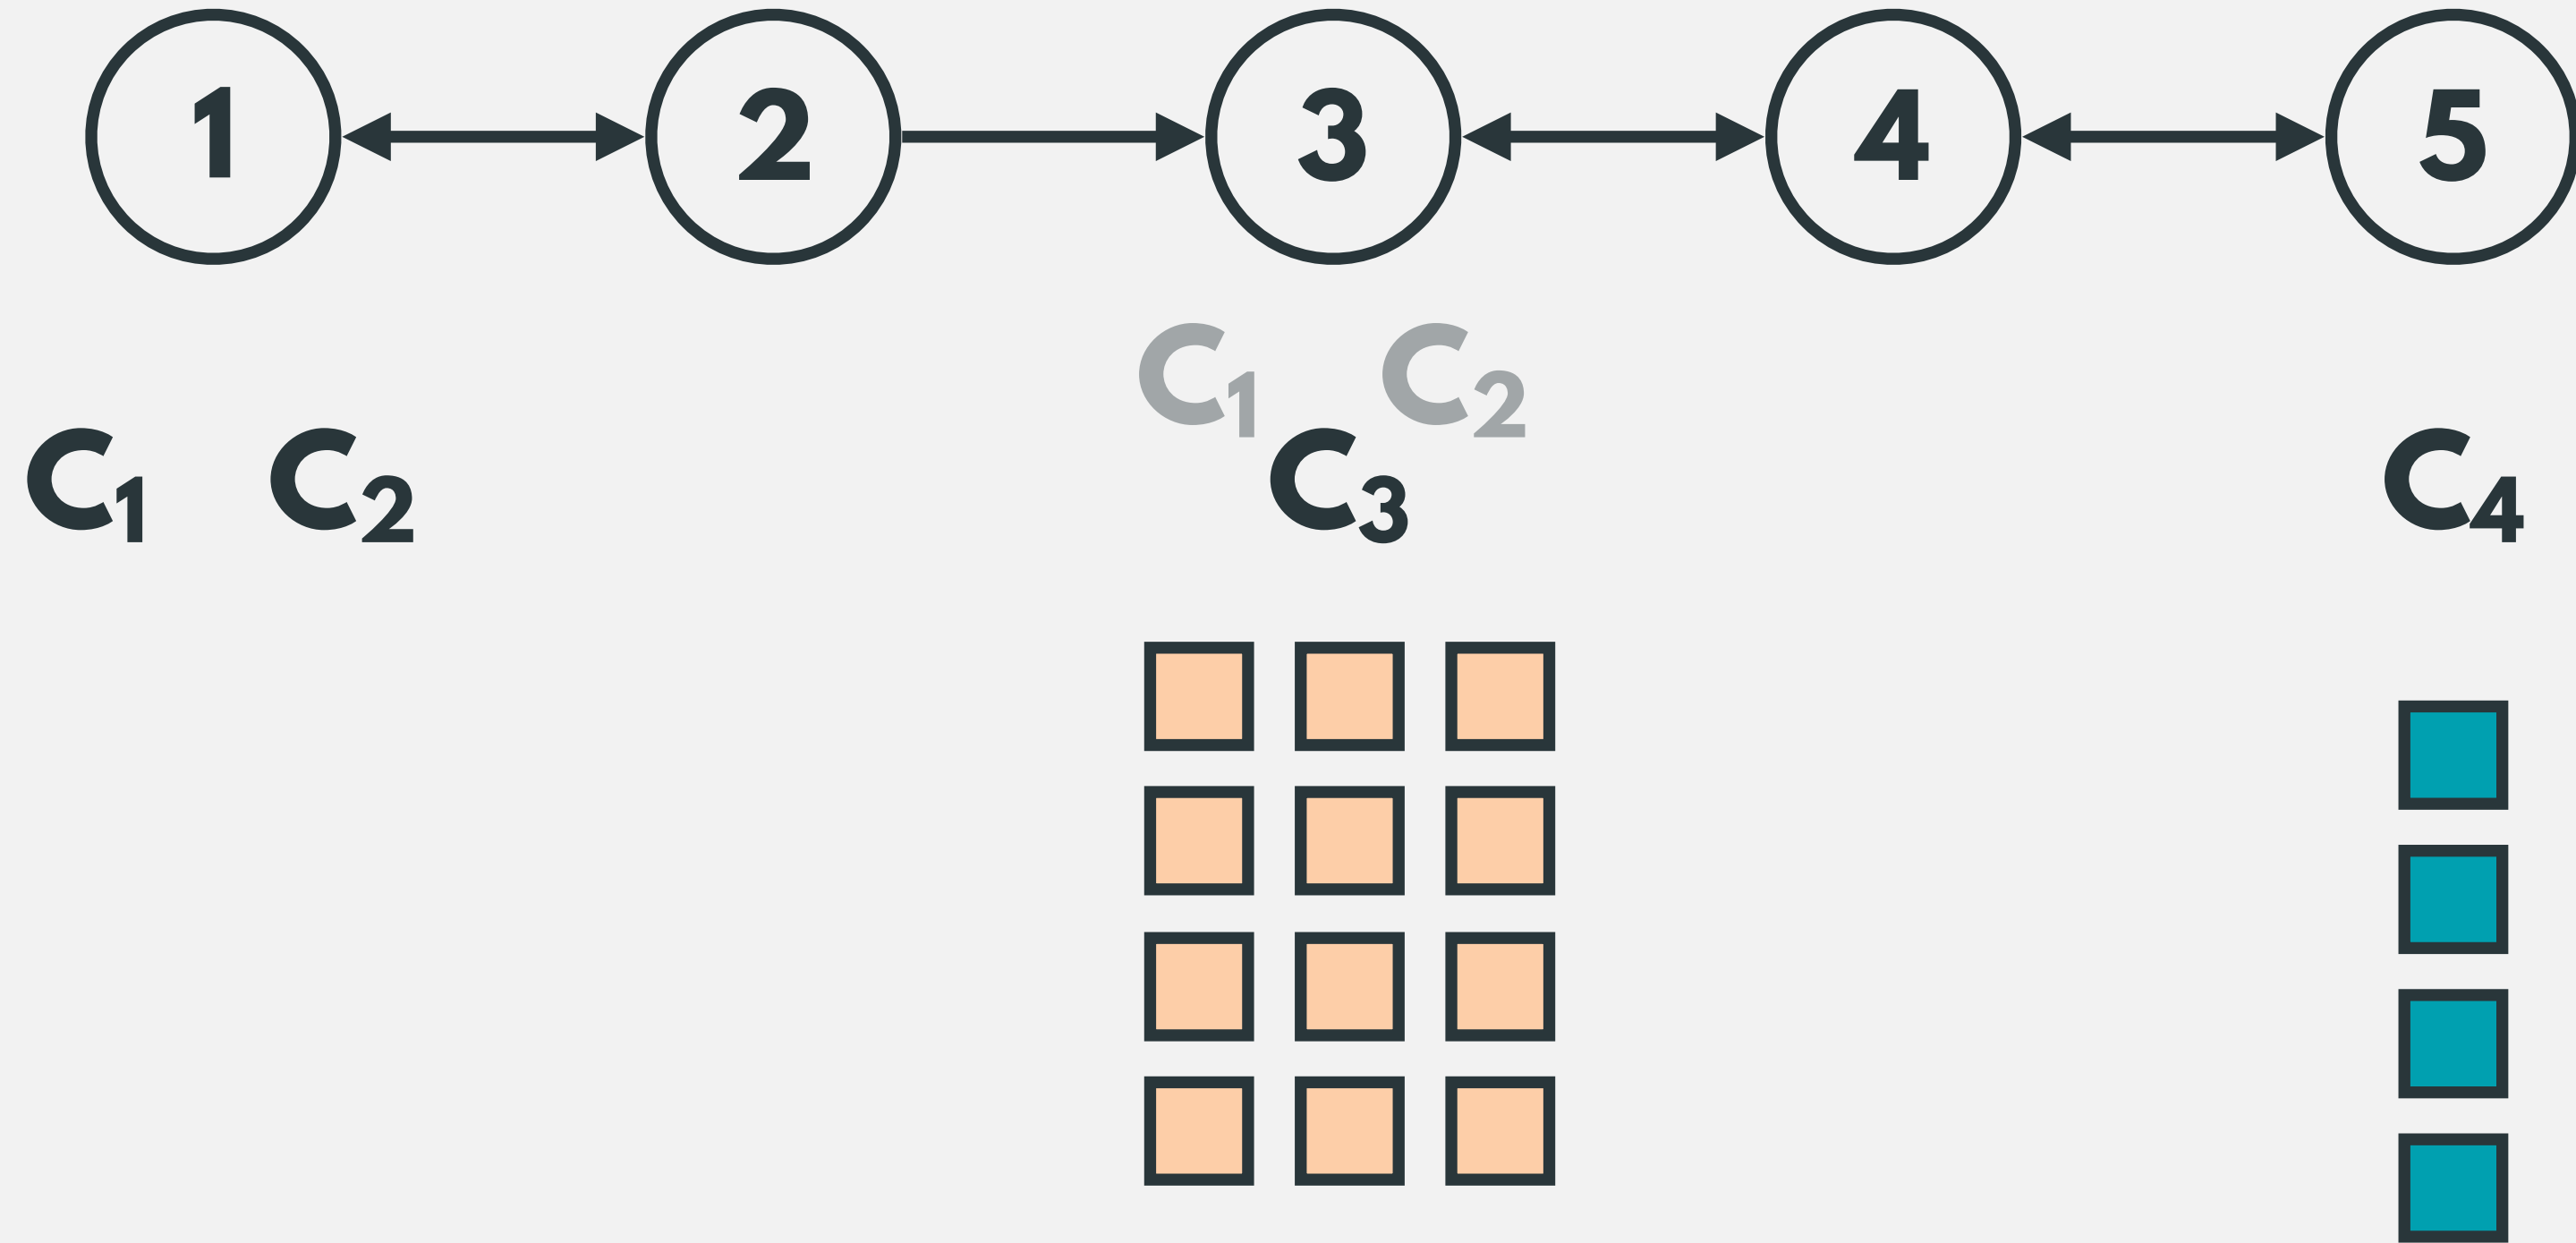
- We'll look at:
 - **Multi-schema** stores (Kafka, et al)
 - **Single-schema** stores (RDBMS, et al)

Multi-Schema Stores






Step 1: Old clients write with new schema



Step 2: Migrate old data to new schema



Step 3: Migrate old clients to new schema

In Practice: Kafka (Confluent)

- Stores Avro schemas in **schema registry** with compatibility check built-in
- Backwards-compatible changes: update **readers** first
- Forwards-compatible changes: update **writers** first


In Practice: Stripe

- **Goal:** (backwards-incompatible) API responses that work for old clients, forever
- Version change modules applied backwards from present to client's version
- They admit: this is **hard**, can't last forever


Recap

- **In multi-schema stores:**
 - **Migrate client writes/reads**
 - **Migrate data**
 - **Migrate client reads/writes**


Single Schema Stores




Single-Schema Migration




**Move from X
to (incompatible) $X + 1$
without downtime**




Step 1: Create and migrate temporary store S'




Step 2: Create a **copier and an **updater****




Step 3.1: Move clients over to new schema



Step 3.2: Copy data, record / apply updates



Step 4: Cutover - S' becomes S



Step 5: Drain updater, delete S_{old}

In Practice - Percona

- pt-online-schema-change
 - **Copier**: scan/copy in timed chunks
 - **Updater**: synchronous table triggers
 - **Cutover**: RENAME TABLE

In Practice - Facebook

- **OSC (Online Schema Change)**
 - **Copier**: dump/load snapshot chunks
 - **Updater**: triggers + changelog
 - **Cutover**: lock, replay, rename, unlock

In Practice - GitHub

- **gh-ost**
 - **Copier**: chunked reads/writes
 - **Updater**: read binlog, interleave copies
 - **Cutover**: 2-step blocking swap

In Practice - Wix

- Updater client-side, controlled with feature toggles
 - **Phase 1:** write to S and S'
 - **Phase 2:** read from S' , fallback to S
 - **Phase 3:** write only to new
- Copy phase eagerly after phase 3

Recap

- **In single-schema stores:**
 - **Migrate clients, retaining old schema**
 - **Create data w/ new schema, over time**
 - **When in sync, cut over**

1. WHAT ARE SCHEMAS?

2. SCHEMA COMPATIBILITY

3. CROSSING COMPATIBILITY GAPS

4. WRAPPING UP / TAKEAWAYS

Wrapping Up

- **Concepts covered here cut across RDBMSs, APIs, document stores, data lakes, etc.**
- **Reason about schema evolution up-front to guide your architecture choices**
- **Your schemas will change. Be prepared.**

A Few Takeaways

- **Make your changes compatible if you can**
- **Compatibility informs migration strategy**
- **Schema-on-read/need won't save you, but it helps sometimes**
- **One size does not fit all**

Thank You! Questions?

Contact Me

Rate This Talk

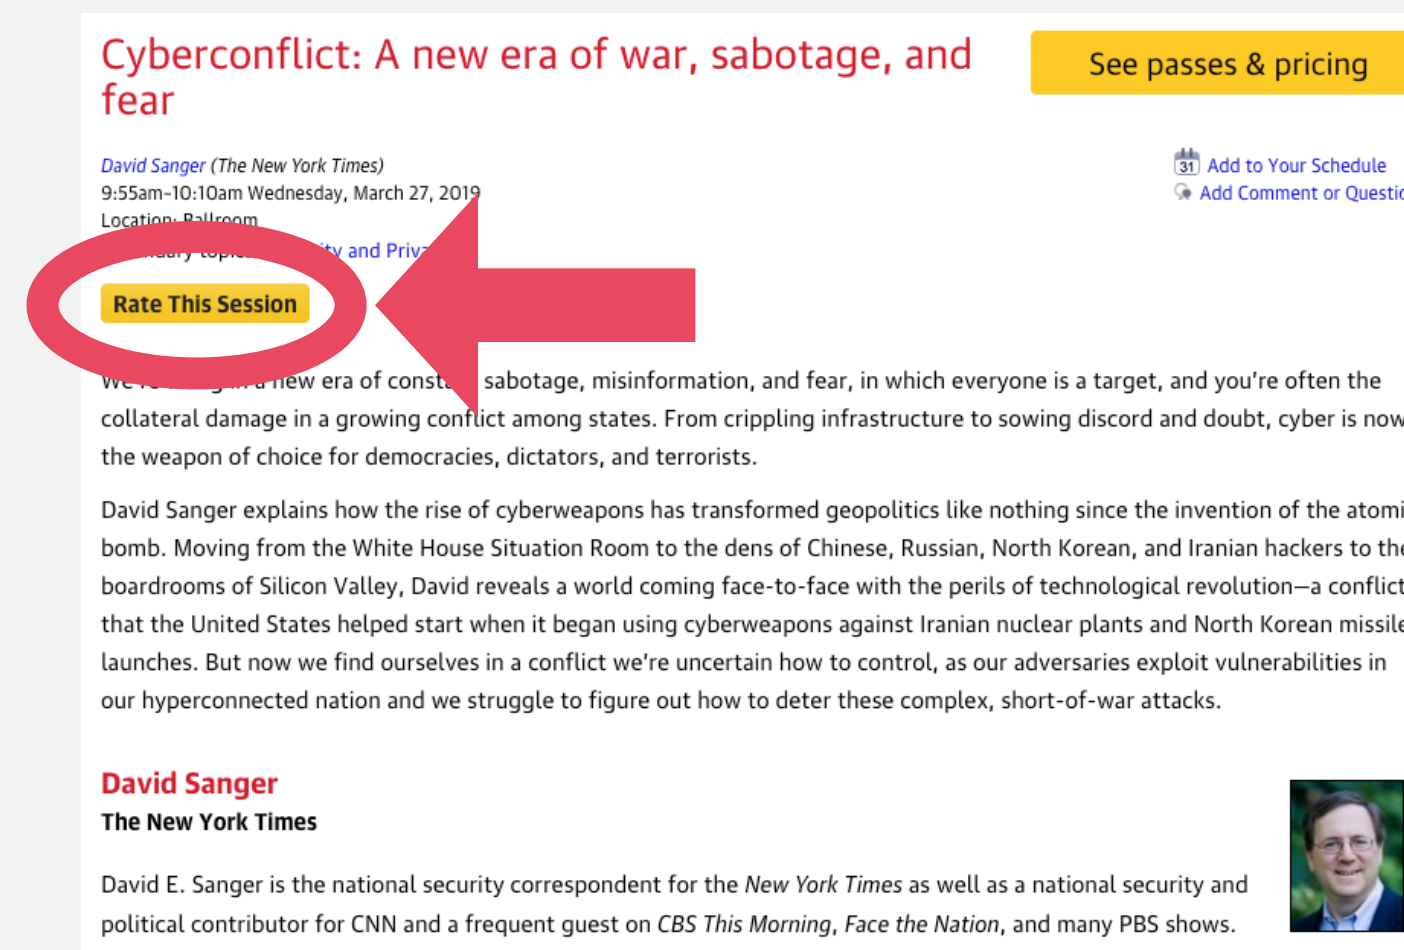


BITS ON DISK

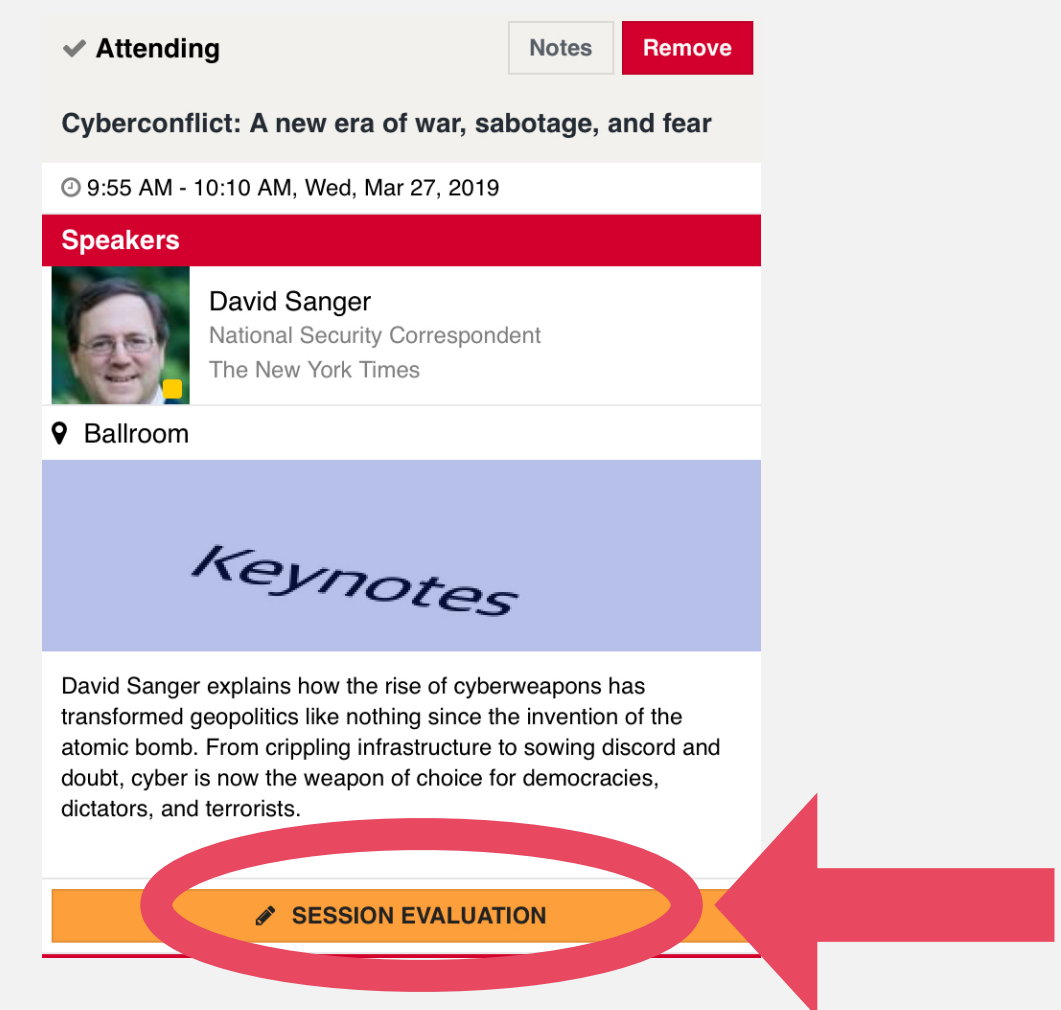
<https://www.bitsondisk.com/>

alex@bitsondisk.com

Twitter, GitHub, etc: @alexras



Session Page



O'Reilly Events App

References: <https://github.com/bitsondisk/velocity-sj-2019>